

## SEATWORK 7.1 Data Wrangling and Notebook Demonstration

**Name:** Calingo, Christian Lei S.

**Section:** CPE22S3

**Course:** Computational Thinking With Python

**Course Code :** CPE311

### 7.5 (HANDLING MISSING DATA)

#### Setup

```
import pandas as pd
df = pd.read_csv('data/dirty_data.csv')
```

#### Finding problematic data

df #as we look in the data, some of them are NaN Values and some are ? Values

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
0	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
1	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
3	2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN	False
4	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
...	...	...	...	...	...	...	...	...	...	...
760	2018-12-31T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	3.3	-3.3	-2.8	NaN	False
761	2018-12-31T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	3.3	-3.3	-2.8	NaN	False
762	2018-12-31T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	3.3	-3.3	-2.8	NaN	False
763	2018-12-31T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
764	2018-12-31T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN

765 rows × 10 columns

[Looking at summary statistics can reveal strange or missing values:](#)

```
df.describe()
```

```
/usr/local/lib/python3.10/dist-packages/numpy/lib/function_base.py:4655: RuntimeWarning: invalid value encountered in subtract
diff_b_a = subtract(b, a)
```

	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	
<b>count</b>	765.000000	577.000000	577.0	765.000000	765.000000	398.000000	11.000000	
<b>mean</b>	5.360392	4.202773	NaN	2649.175294	-15.914379	8.632161	16.290909	
<b>std</b>	10.002138	25.086077	NaN	2744.156281	24.242849	9.815054	9.489832	
<b>min</b>	0.000000	0.000000	-inf	-11.700000	-40.000000	-16.100000	1.800000	
<b>25%</b>	0.000000	0.000000	NaN	13.300000	-40.000000	0.150000	8.600000	
<b>50%</b>	0.000000	0.000000	NaN	32.800000	-11.100000	8.300000	19.300000	
<b>75%</b>	5.800000	0.000000	NaN	5505.000000	6.700000	18.300000	24.900000	
<b>max</b>	61.700000	229.000000	inf	5505.000000	23.900000	26.100000	28.700000	

[Using info\(\) function, we can pinpoint the missing values and incorrect data types](#)

```
df.info() # The number under the non-null count shows us how many rows doesn't have NaN values with accordance of its column
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 765 entries, 0 to 764
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  765 non-null   object
1   station               765 non-null   object
2   PRCP                  765 non-null   float64
3   SNOW                  577 non-null   float64
4   SNWD                  577 non-null   float64
5   TMAX                  765 non-null   float64
6   TMIN                  765 non-null   float64
7   TOBS                  398 non-null   float64
8   WESF                  11 non-null    float64
9   inclement_weather    408 non-null   object
dtypes: float64(7), object(3)
memory usage: 59.9+ KB
```

[With the use of pd.isnull\(\) / pd.isna\(\) or the isna\(\) / isnull\(\), we can find nulls](#)

```
contain_nulls = df[
df.SNOW.isnull() | df.SNWD.isna()\
| pd.isnull(df.TOBS) | pd.isna(df.WESF)\
| df.inclement_weather.isna()
]
contain_nulls.shape

# if we remove the index in .shape(), we can see a tuple
# the index[0] is the number of nulls while the index[1] is the number of columns

(765, 10)
```

```
contain_nulls.head(10)
```

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
<b>0</b>	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
<b>1</b>	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
<b>2</b>	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
<b>3</b>	2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN	False
<b>4</b>	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
<b>5</b>	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
<b>6</b>	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
<b>7</b>	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
<b>8</b>	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
<b>9</b>	2018-01-05T00:00:00	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	NaN

[We can't check if we have NaN like this](#)

```
df[df.inclement_weather == 'NaN'].shape[0]

0
```

[This is because it is actually np.nan . However, notice this also doesn't work](#)

```
import numpy as np
df[df.inclement_weather == np.nan].shape[0]

0
```

[We have to use one of the methods discussed earlier for this to work](#)

```
df[df.inclement_weather.isna()].shape # inclement_weather have 357 NaN Values while the
                                     # 10 represents the number of columns in the data

(357, 10)
```

[We can find -inf / inf by comparing to -np.inf / np.inf](#)

```
df[df.SNWD.isin([-np.inf, np.inf])).shape[0]

577
```



[Rather than do this for each column, we can write a function that will use a dictionary comprehension to check all the columns for us:](#)

```
import numpy as np
def get_inf_count(df):
    return {
        col : df[df[col].isin([np.inf, -np.inf])).shape[0] for col in df.columns
    }
get_inf_count(df)

{'date': 0,
 'station': 0,
 'PRCP': 0,
 'SNOW': 0,
 'SNWD': 577,
 'TMAX': 0,
 'TMIN': 0,
 'TOBS': 0,
 'WESF': 0,
 'inclement_weather': 0}
```

[Before we can decide how to handle the infinite values of snow depth, we should look at the summary statistics for snowfall which form a big part in determining the snow depth](#)

```
pd.DataFrame({
'np.inf Snow Depth': df[df.SNWD == np.inf].SNOW.describe(),
'-np.inf Snow Depth': df[df.SNWD == -np.inf].SNOW.describe()
}).T
```

	count	mean	std	min	25%	50%	75%	max	
np.inf Snow Depth	24.0	101.041667	74.498018	13.0	25.0	120.5	152.0	229.0	
-np.inf Snow Depth	553.0	0.000000	0.000000	0.0	0.0	0.0	0.0	0.0	

[Let's now look into the date and station columns. We saw the ? for station earlier, so we know that was the other unique value. However, we see that some dates are present 8 times in the data and we only have 324 days meaning we are also missing days](#)

```
df.describe(include='object')
```

	date	station	inclement_weather
count	765	765	408
unique	324	2	2
top	2018-07-05T00:00:00	GHCND:USC00280907	False
freq	8	398	384

[We can use the duplicated\(\) method to find duplicate rows](#)

```
df[df.duplicated()].shape # same like the above, the first index shows the number of duplicated values in the data
# while the second index shows the number of columns present in the data
```

```
(284, 10)
```

[The default for keep is 'first' meaning it won't show the first row that the duplicated data was seen in; we can pass in False to see it though:](#)

```
df[df.duplicated(keep=False)].shape[0]
```

```
482
```

[We can also specify the columns to use:](#)

```
df[df.duplicated(['date', 'station'])].shape[0]
```

```
284
```

**Trying other columns**

```
df[df.duplicated(['inclement_weather', 'PRCP'])].shape
```

```
(636, 10)
```

[Let's look at a few duplicates. Just in the few values we see here, we know that the top 4 are actually in the data 6 times because by default we aren't seeing their first occurrence:](#)

```
df[df.duplicated()].head()
```

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
1	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
5	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
6	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
8	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True

## ✓ Mitigating Issues

### Handling duplicated data

[Since we know we have NY weather data and noticed we only had two entries for station, we may decide to drop the station column because we are only interested in the weather data. However, when dealing with duplicate data, we need to think of the ramifications of removing it. Notice we only have data for the WESF column when the station is ? :](#)

```
df[df.WESF.notna()].station.unique()
```

```
array(['?'], dtype=object)
```

```
station_qm_wesf = df[df.station == '?'].WESF

df.sort_values('station', ascending=False, inplace=True)

df_deduped = df.drop_duplicates('date').drop(

columns=['station', 'WESF']

).sort_values('date').assign(

WESF=station_qm_wesf
)
df_deduped.shape

(324, 9)
```

[Check out the 4th row, we have WESF in the correct spot thanks to the index](#)

```
df_deduped.head()
```

		date	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF	
1		2018-01-01T00:00:00	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN	
3		2018-01-02T00:00:00	0.0	0.0	-inf	-8.3	-16.1	-12.2	False	NaN	
6		2018-01-03T00:00:00	0.0	0.0	-inf	-4.4	-13.9	-13.3	False	NaN	
8		2018-01-04T00:00:00	20.6	229.0	inf	5505.0	-40.0	NaN	True	19.3	
11		2018-01-05T00:00:00	14.2	127.0	inf	-4.4	-13.9	-13.9	True	NaN	

## ✓ Dealing with nulls

```
df_deduped.dropna().shape

(0, 9)
```

[Using how = all, we can choose to only drop the rows where everything is null](#)

```
df_deduped.dropna(how='all').shape

# This doesn't drop anything since there are no rows where everything is null

(324, 9)
```

[We can use just a subset of columns to determine what to drop with the subset argument](#)

```
df_deduped.dropna(
how='all', subset=['inclement_weather', 'SNOW', 'SNWD'] # in this case, we are dropping rows with NaN values with the following columns
).shape

(293, 9)
```

### Trying Other Columns

```
df_deduped.dropna(
how='all', subset=['TOBS', 'WESF']
).shape

(251, 9)
```

[This can also be performed along columns, and we can also require a certain number of null values before we drop the data](#)

```
df_deduped.dropna(axis='columns', thresh=df_deduped.shape[0]*.75).columns
```

```
Index(['date', 'PRCP', 'SNOW', 'SNWD', 'TMAX', 'TMIN', 'TOBS',
      'inclement_weather'],
      dtype='object')
```

We can choose to fill in the null values instead with fillna() [link text](#)

```
df_deduped.loc[:, 'WESF'].fillna(0, inplace=True) # we are locating the WESF column and filling the NaN values with 0
df_deduped.head()
```

	date	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF
1	2018-01-01T00:00:00	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	0.0
3	2018-01-02T00:00:00	0.0	0.0	-inf	-8.3	-16.1	-12.2	False	0.0
6	2018-01-03T00:00:00	0.0	0.0	-inf	-4.4	-13.9	-13.3	False	0.0
8	2018-01-04T00:00:00	20.6	229.0	inf	5505.0	-40.0	NaN	True	19.3
11	2018-01-05T00:00:00	14.2	127.0	inf	-4.4	-13.9	-13.9	True	0.0

Trying to fill other columns with another value

```
df_deduped.loc[:, 'TOBS'].fillna(-15, inplace=True)
df_deduped.head()
```

#the NaN values in the TOBS column is now replaced with -15

	date	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF
1	2018-01-01T00:00:00	0.0	0.0	-inf	5505.0	-40.0	-15.0	NaN	0.0
3	2018-01-02T00:00:00	0.0	0.0	-inf	-8.3	-16.1	-12.2	False	0.0
6	2018-01-03T00:00:00	0.0	0.0	-inf	-4.4	-13.9	-13.3	False	0.0
8	2018-01-04T00:00:00	20.6	229.0	inf	5505.0	-40.0	-15.0	True	19.3
11	2018-01-05T00:00:00	14.2	127.0	inf	-4.4	-13.9	-13.9	True	0.0

At this point we have done every we can without distorting the data. We know that we are missing dates, but if we reindex, we don't know how to fill in the NaN data. With the weather data, we can't assume because it snowed one day that it will snow the next or that the temperature will be the same. For this reason, note that the next few examples are just for illustrative purposes only—just because we can do something doesn't mean we should. That being said, let's try to address some of remaining issues with the temperature data. We know that when TMAX is the temperature of the Sun, it must be because there was no measured value, so let's replace it with NaN and then we will make an assumption that the temperature won't change drastically day-to-day. Note that this is actually a big assumption, but it will allow us to understand how fillna() works when we provide a strategy through the method parameter. We will also do this for TMIN which currently uses -40°C for its placeholder when we know that the coldest temperature ever recorded in NYC was -15°F (-26.1°C) on February 9, 1934. In [23]: df\_deduped.dropna(how='all', subset=['inclement\_weather', 'SNOW', 'SNWD']).shape Out[23]: In [24]: df\_deduped.dropna(axis='columns', thresh=df\_deduped.shape[0]\*.75).columns Out[24]: In [25]: df\_deduped.loc[:, 'WESF'].fillna(0, inplace=True) df\_deduped.head() Out[25]:

The fillna() method gives us 2 options for the method parameter:

- 'ffill' to forward fill
- 'bfill' to back fill

```
df_deduped.assign(
    TMAX=lambda x: x.TMAX.replace(5505, np.nan).fillna(method='ffill'),
    TMIN=lambda x: x.TMIN.replace(-40, np.nan).fillna(method='ffill')
).head()
```

# The TOBS Column did not become NaN since we have tried to fill it with non NaN values in the previous cell

	date	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF	
1	2018-01-01T00:00:00	0.0	0.0	-inf	NaN	NaN	-15.0		NaN	0.0
3	2018-01-02T00:00:00	0.0	0.0	-inf	-8.3	-16.1	-12.2		False	0.0
6	2018-01-03T00:00:00	0.0	0.0	-inf	-4.4	-13.9	-13.3		False	0.0
8	2018-01-04T00:00:00	20.6	229.0	inf	-4.4	-13.9	-15.0		True	19.3
11	2018-01-05T00:00:00	14.2	127.0	inf	-4.4	-13.9	-13.9		True	0.0

[We can use np.nan\\_to\\_num\(\) to turn np.nan into 0 and -np.inf / np.inf into large negative or positive finite numbers](#)

```
df_deduped.assign(
    SNWD=lambda x: np.nan_to_num(x.SNWD)
).head()
```

	date	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF	
1	2018-01-01T00:00:00	0.0	0.0	-1.797693e+308	5505.0	-40.0	-15.0		NaN	0.0
3	2018-01-02T00:00:00	0.0	0.0	-1.797693e+308	-8.3	-16.1	-12.2		False	0.0
6	2018-01-03T00:00:00	0.0	0.0	-1.797693e+308	-4.4	-13.9	-13.3		False	0.0
8	2018-01-04T00:00:00	20.6	229.0	1.797693e+308	5505.0	-40.0	-15.0		True	19.3
11	2018-01-05T00:00:00	14.2	127.0	1.797693e+308	-4.4	-13.9	-13.9		True	0.0

[We can couple fillna\(\) with other types of calculations for interpolation. Here we replace missing values of TMAX with the median of all TMAX values, TMIN with the median of all TMIN values, and TOBS to the average of the TMAX and TMIN values. Since we place TOBS last, we have access to the imputed values for TMIN and TMAX in the calculation. WARNING: the text has a typo and fills in TMAX with TMIN's median, the below is correct.](#)

```
df_deduped.assign(
    TMAX=lambda x: x.TMAX.replace(5505, np.nan).fillna(x.TMAX.median()),
    TMIN=lambda x: x.TMIN.replace(-40, np.nan).fillna(x.TMIN.median()),
    # average of TMAX and TMIN
    TOBS=lambda x: x.TOBS.fillna((x.TMAX + x.TMIN) / 2)
).head()
```

	date	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF	
1	2018-01-01T00:00:00	0.0	0.0	-inf	22.8	0.0	-15.0		NaN	0.0
3	2018-01-02T00:00:00	0.0	0.0	-inf	-8.3	-16.1	-12.2		False	0.0
6	2018-01-03T00:00:00	0.0	0.0	-inf	-4.4	-13.9	-13.3		False	0.0
8	2018-01-04T00:00:00	20.6	229.0	inf	22.8	0.0	-15.0		True	19.3
11	2018-01-05T00:00:00	14.2	127.0	inf	-4.4	-13.9	-13.9		True	0.0

[We can also use apply\(\) for running the same calculation across columns. For example, let's fill all missing values with their rolling 7 day median of their values, setting the number of periods required for the calculation to 0 to ensure we don't introduce more extra NaN values. \(Rolling calculations will be covered in chapter 4.\) We need to set the date column as the index so apply\(\) doesn't try to take the rolling 7 day median of the date](#)

```
df_deduped.assign(

    TMAX=lambda x: x.TMAX.replace(5505, np.nan),
    TMIN=lambda x: x.TMIN.replace(-40, np.nan)
).set_index('date').apply(

    lambda x: x.fillna(x.rolling(7, min_periods=0).median())
).head(10)
```

	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF
date								
2018-01-01T00:00:00	0.0	0.0	-inf	NaN	NaN	-15.0	NaN	0.0
2018-01-02T00:00:00	0.0	0.0	-inf	-8.30	-16.1	-12.2	False	0.0
2018-01-03T00:00:00	0.0	0.0	-inf	-4.40	-13.9	-13.3	False	0.0
2018-01-04T00:00:00	20.6	229.0	inf	-6.35	-15.0	-15.0	True	19.3
2018-01-05T00:00:00	14.2	127.0	inf	-4.40	-13.9	-13.9	True	0.0
2018-01-06T00:00:00	0.0	0.0	-inf	-10.00	-15.6	-15.0	False	0.0
2018-01-07T00:00:00	0.0	0.0	-inf	-11.70	-17.2	-16.1	False	0.0
2018-01-08T00:00:00	0.0	0.0	-inf	-7.80	-16.7	-8.3	False	0.0

[The last strategy we could try is interpolation with the interpolate\(\) method. We specify the method parameter with the interpolation strategy to use. There are many options, but we will stick with the default of 'linear', which will treat values as evenly spaced and place missing values in the middle of existing ones. We have some missing data, so we will reindex first. Look at January 9th, which we didn't have before—the values for TMAX, TMIN, and TOBS are the average of values the day prior \(January 8th\) and the day after \(January 10th\).](#)

```
df_deduped.assign(
    TMAX=lambda x: x.TMAX.replace(5505, np.nan),
    TMIN=lambda x: x.TMIN.replace(-40, np.nan),
    date=lambda x: pd.to_datetime(x.date)
).set_index('date').reindex(
    pd.date_range('2018-01-01', '2018-12-31', freq='D')
).apply(
    lambda x: x.interpolate()
).head(10)
```

	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	inclement_weather	WESF
date								
2018-01-01	0.0	0.0	-inf	NaN	NaN	-15.00	NaN	0.0
2018-01-02	0.0	0.0	-inf	-8.3	-16.10	-12.20	False	0.0
2018-01-03	0.0	0.0	-inf	-4.4	-13.90	-13.30	False	0.0
2018-01-04	20.6	229.0	inf	-4.4	-13.90	-15.00	True	19.3
2018-01-05	14.2	127.0	inf	-4.4	-13.90	-13.90	True	0.0
2018-01-06	0.0	0.0	-inf	-10.0	-15.60	-15.00	False	0.0
2018-01-07	0.0	0.0	-inf	-11.7	-17.20	-16.10	False	0.0